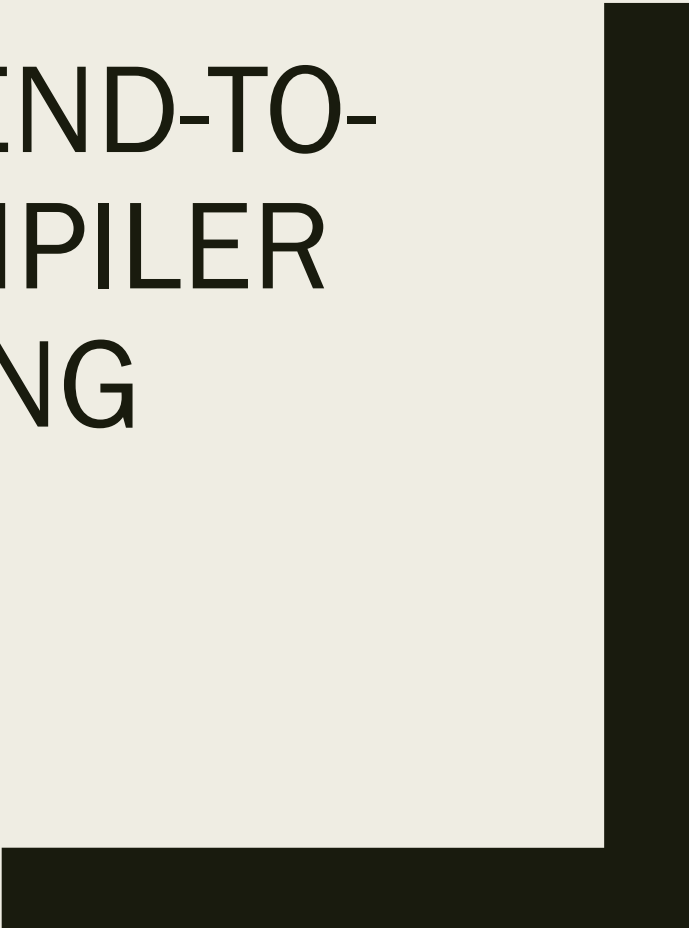




TVM: AN AUTOMATED END-TO- END OPTIMIZING COMPILER FOR DEEP LEARNING

전기정보공학부 김희원

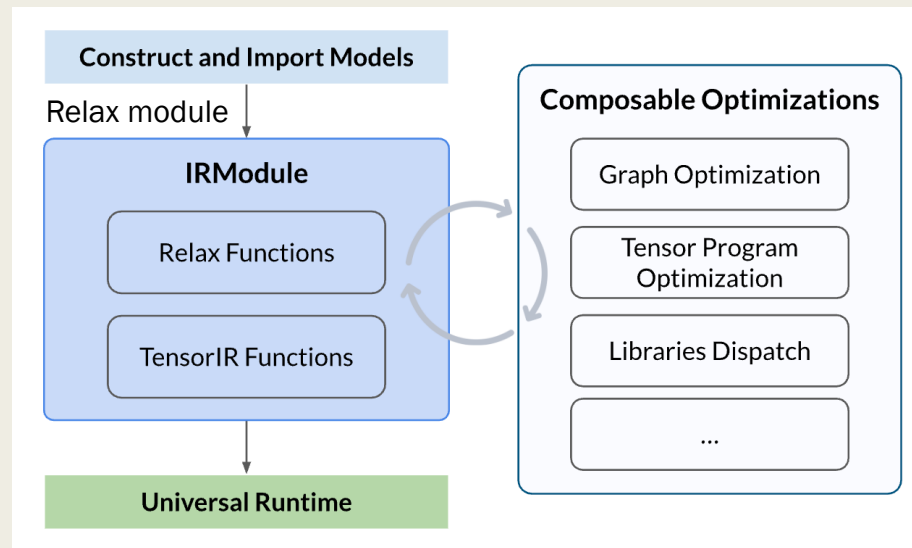


Papers read

- Chen, T., Moreau, T., Jiang, Z., Zheng, L., Yan, E., Shen, H., Cowan, M., Wang, L., Hu, Y., Ceze, L., Guestrin, C., & Krishnamurthy, A. (2018). *TVM: An automated end-to-end optimizing compiler for deep learning*. In *Proceedings of the 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI '18)* (pp. 579–594). USENIX Association.
<https://www.usenix.org/conference/osdi18/presentation/chen>
- Chen, T., Zheng, L., Yan, E., Jiang, Z., Moreau, T., Ceze, L., Guestrin, C., & Krishnamurthy, A. (2018). *Learning to optimize tensor programs*. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, & R. Garnett (Eds.), *Advances in Neural Information Processing Systems* (Vol. 31, pp. 3393–3404). Curran Associates, Inc. <https://doi.org/10.48550/arXiv.1805.08166>

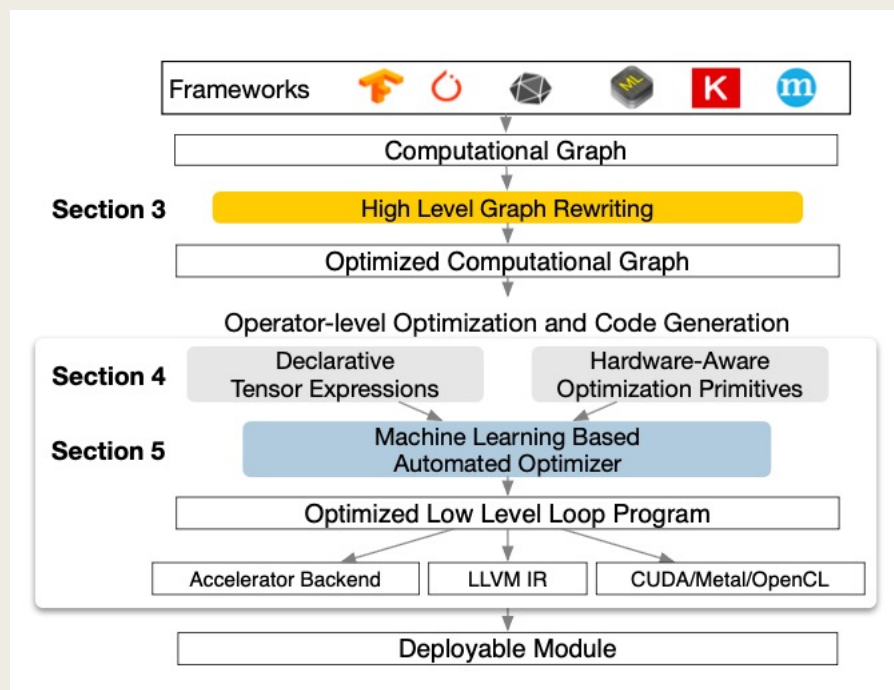
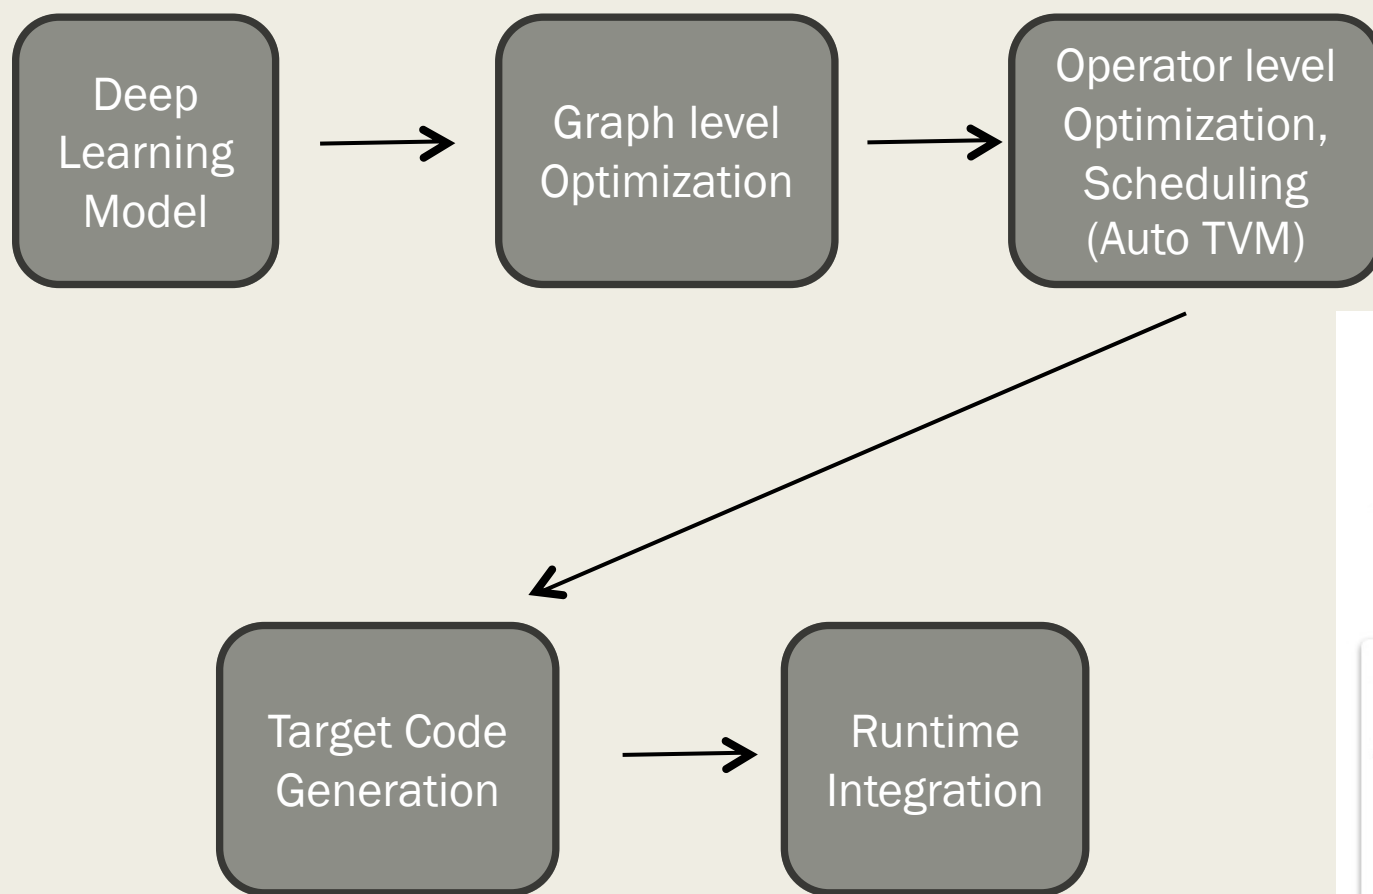
Overview

- End-to-end deep learning compiler stack
- Translates high-level models to low level optimized code
- Supports diverse backends: CPU, GPU, FPGA
- Graph-level optimization, operator-level optimization: auto-scheduling via AutoTVM



https://tvm.apache.org/docs/how_to/tutorials/e2e_opt_model.html

TVM Compiler Pipeline




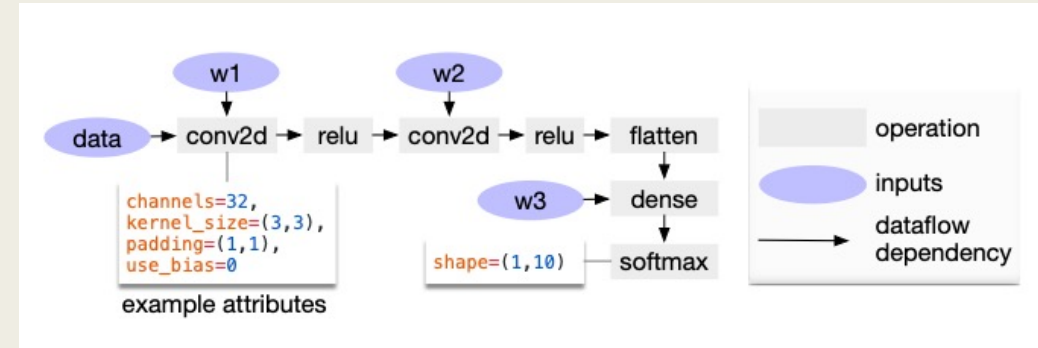
1. Graph-Level Optimizations

- **Computational Graph**

Node = operation / Edge = dataflow dependency

- **Graph-level optimization methods**

- **Operator Fusion:** combine multiple operators into a single kernel
- **Constant Folding:** precompute statically determinable values at compile-time
- **Static Memory Planning:** pre-allocate memory for each intermediate tensor
- **Data Layout Transformation:** reorder data in back-end-friendly forms for memory efficiency



2. Operator-Level Optimizations

- Tensor expression language for $C = A^T \cdot B$:

```
m, n, h = t.var('m'), t.var('n'), t.var('h')
A = t.placeholder((m, h), name='A')
B = t.placeholder((n, h), name='B')
k = t.reduce_axis((0, h), name='k')
C = t.compute((m, n), lambda y, x:
               t.sum(A[k, y] * B[k, x], axis=k))
```

computing rule

result shape

- Doesn't specify execution details such as loop structure- separates scheduling from computation
- **1) Nested parallelism with cooperation:** data parallel tasks are parallelized.
“Shared-nothing nested parallelism”: not shared between working thread siblings.

```
for thread_group (by, bx) in cross(64, 64):
    for thread_item (ty, tx) in cross(2, 2):
        local CL[8][8] = 0
        shared AS[2][8], BS[2][8]
        for k in range(1024):
            for i in range(4):
                AS[ty][i*4+tx] = A[k][by*64+ty*8+i*4+tx]
            for each i in 0..4:
                BS[ty][i*4+tx] = B[k][bx*64+ty*8+i*4+tx]
            memory_barrier_among_threads()
            for yi in range(8):
                for xi in range(8):
                    CL[yi][xi] += AS[yi] * BS[xi]
            for yi in range(8):
                for xi in range(8):
                    C[yo*8+yi][xo*8+xi] = CL[yi][xi]
```

All threads cooperatively
load AS and BS in different
parallel patterns

Barrier inserted
automatically
by compiler

AS, BS: shared arrays

2. Operator-Level Optimizations

■ 2) Tensorization:

- Computations → Tensor operations
- *TVM made tensorization extensible by separating the target hardware intrinsic (instructions) from the schedule.*

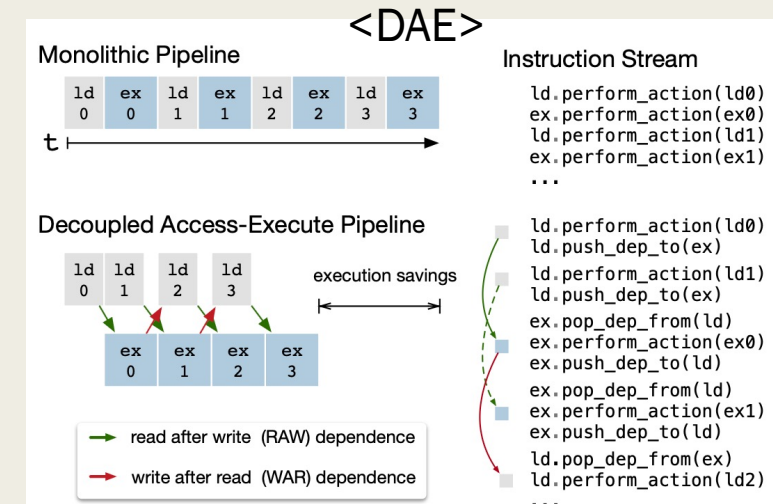
```
w, x = t.placeholder((8, 8)), t.placeholder((8, 8))
k = t.reduce_axis((0, 8))
y = t.compute((8, 8), lambda i, j:
    t.sum(w[i, k] * x[j, k], axis=k))
def gemm_intrin_lower(inputs, outputs):
    ww_ptr = inputs[0].access_ptr("r")
    xx_ptr = inputs[1].access_ptr("r")
    zz_ptr = outputs[0].access_ptr("w")
    compute = t.hardware_intrin("gemm8x8", ww_ptr, xx_ptr, zz_ptr)
    reset = t.hardware_intrin("fill_zero", zz_ptr)
    update = t.hardware_intrin("fuse_gemm8x8_add", ww_ptr, xx_ptr, zz_ptr)
    return compute, reset, update
gemm8x8 = t.decl_tensor_intrin(y.op, gemm_intrin_lower)
```

declare behavior

lowering rule to generate hardware intrinsics to carry out the computation

■ 3) Explicit Memory Latency Hiding: The process of overlapping memory operations with computation to maximize utilization of memory and compute resources.

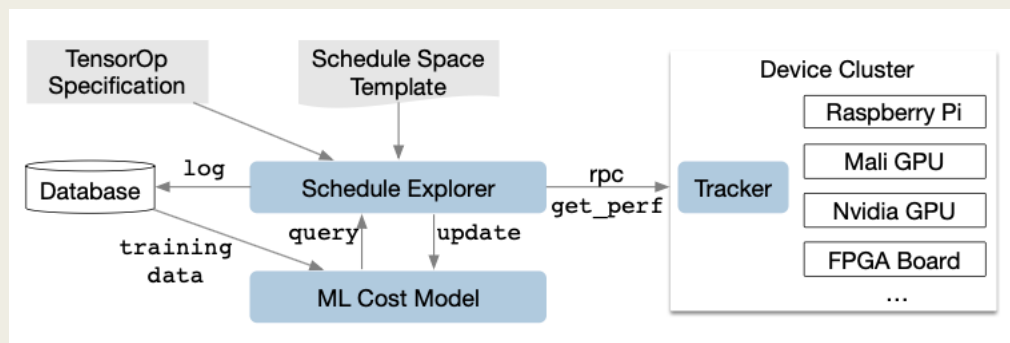
- **CPU**: simultaneous multithreading or hardware prefetching
- **GPU**: rapid context switching of many warps of threads
- **TPU**: decoupled access-execute (DAE) architecture



2. Operator-Level Optimizations- Auto Scheduling via AutoTVM

■ Motivation:

- Systems relied on manually optimized libraries (ex) cuDNN) → significant engineering cost was needed to deploy them to new hardware targets.
- Use learning to automatically optimize operators.



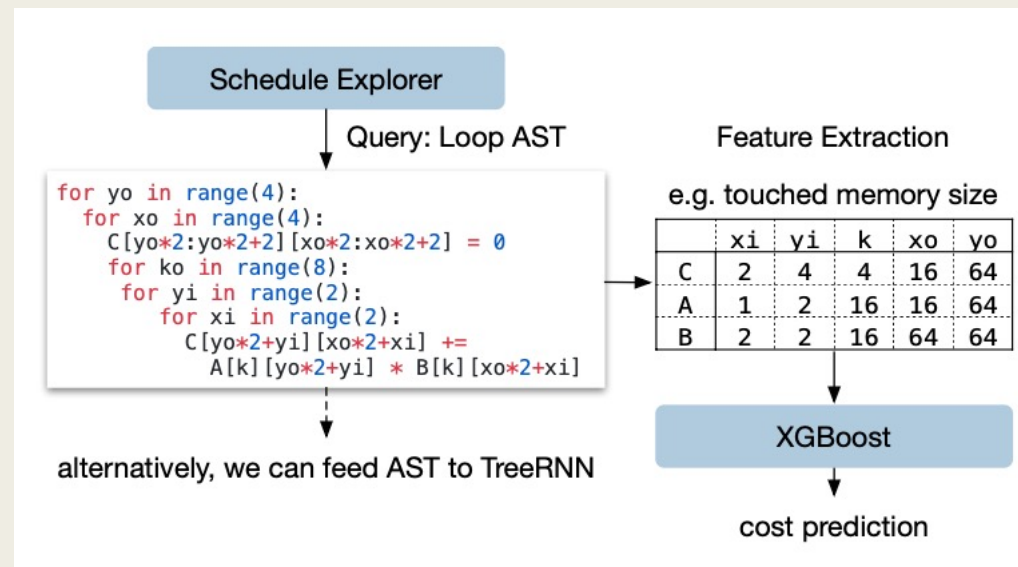

■ 1) Schedule Space Specification

- Considers as many configurations as possible (billions)
- Developers can declare knobs themselves, or, a generic master template for each hardware back-end is created beforehand

2. Operator-Level Optimizations- Auto Scheduling via AutoTVM

■ 2) ML-Based Cost Model

- Statistical approach instead of considering all factors affecting performance
- The model is updated experimentally as we explore more configurations.
- The gradient tree boosting model (based on XGBoost) does prediction in 0.67ms (on average), which was the fastest among different automation methods for a conv2d operator in ResNet-18 on TITAN X.



2. Operator-Level Optimizations- Auto Scheduling via AutoTVM

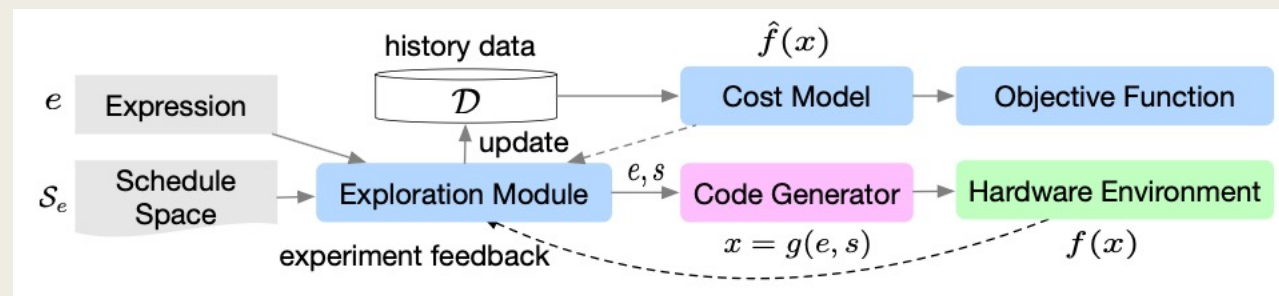
■ 3) Schedule Exploration

- Runs a parallel Simulated Annealing (SA) algorithm for large search spaces.
 - Randomly walks to a nearby configuration:
 - Accepts if cost decreases.
 - Rejects with some probability if cost increases.

■ 4) Distributed Device Pool

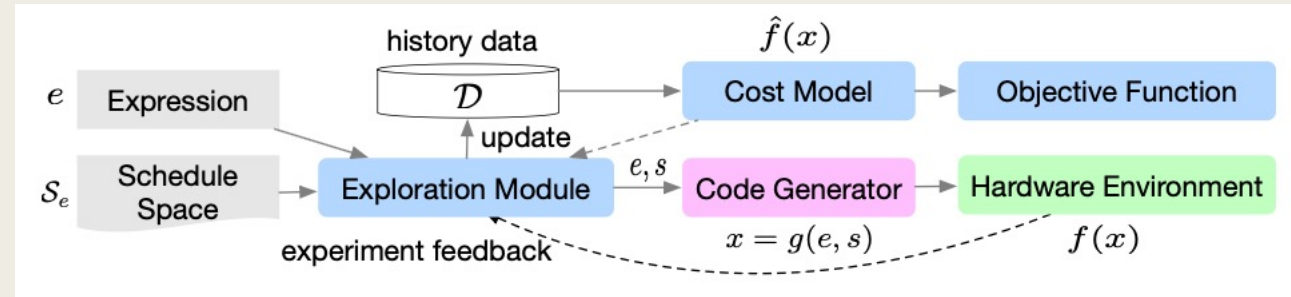
- Implements an RPC-based distributed device pools for optimization.

AutoTVM in detail



- 1) Statistical Cost Model ($\hat{f}(x) \approx f(x)$ for a low-level program x)
 - Extract features
 - Loop Structure Information: memory access count, data reuse ratio...
 - Generic Annotations: vectorization, unrolling, thread binding...
 - i) **XGBoost**: requires precise feature engineering + fast prediction
 - ii) **TreeGRU**: feature engineering not required + training speed is slow
→ used with batching and GPU
- 2) Training Objective Function
 - Run time statistics data $D = \{e_i, s_i, c_i\}$
 - Rank loss function $\sum_{i,j} \log(1 + e^{-\text{sign}(c_i - c_j)(\hat{f}(x_i) - \hat{f}(x_j))})$.
Used because relative order of program run time is important

AutoTVM in detail



■ 3) Exploration Module

- Select the top-performing batch of candidates to update \hat{f} .
- Simulated Annealing

■ 4) Accelerating Optimization via Transfer Learning

- The cost model is shared using the common transferable representation (invariant to the source/target domains)
- Transferable representation:
 - GBT: Z : context feature matrix such that $Z_{ik} = i - th \text{ feature of loop } k$
 Relation feature between features i, j : $R_t^{(ij)} = \max_{k \in \{k | Z_{kj} < \beta_t\}} Z_{ki}$
 - TreeGRU: $\hat{f}(x) = \hat{f}^{(global)}(x) + \hat{f}^{(local)}(x)$
 $\hat{f}^{(global)}$ is trained on D' (= historical data from previously seen workloads)

Comparison to Other Systems

- XLA:
 - Focuses on graph-level optimizations: too high-level
- Halide:
 - Introduced the idea of separating computing and scheduling
→ TVM adopted this + focuses on new scheduling challenges of DL workloads
- ATLAS, FFTW:
 - Autotuning (Single kernel, predefined parameters)
- Tensor Comprehension:
 - Black-box auto-tuning + polyhedral optimizations via Genetic Algorithm

Evaluation

- Evaluated TVM on 4 types of platforms:
 - *1) server-class GPU, 2) embedded GPU, 3) embedded CPU, 4) FPGA SoC*
- Benchmarks:
 - *ResNet, MobileNet, LSTM Language Model, Deep Q Network(DQN), Deep Convolutional Generative Adversarial Networks(DCGAN)*
- Compared with:
 - *Existing DL frameworks such as MxNet, TensorFlow*

Evaluation

■ 1) Server-class GPU (Nvidia Titan X):

- End-to-end performance:
TVM outperforms, due to joint graph optimization + Auto TVM
- Operator-level optimization performance:
TVM mostly outperforms, due to large schedule space + ML-based search

■ 2) Embedded CPU (ARM Cortex A53, Quad Core):

- End-to-end performance:
TVM outperforms
- Operator-level optimization performance:
TVM outperforms hand-optimized versions

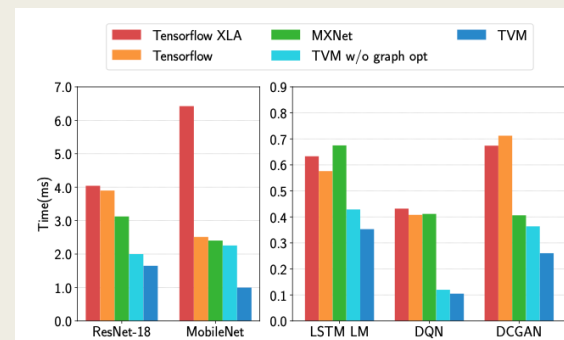


Figure 14: GPU end-to-end evaluation for TVM, MXNet, Tensorflow, and Tensorflow XLA. Tested on the NVIDIA Titan X.

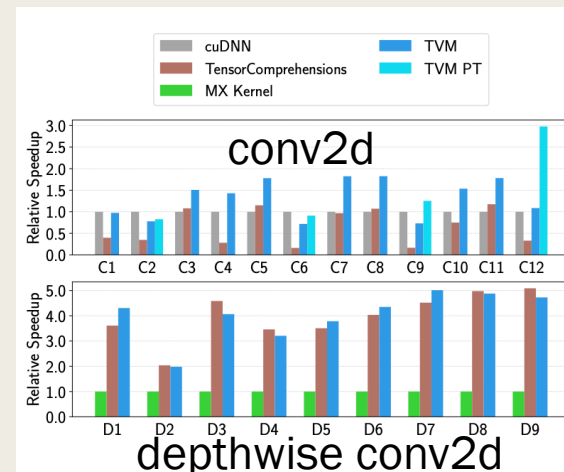


Figure 15: Relative speedup of all conv2d operators in ResNet-18 and all depthwise conv2d operators in MobileNet. Tested on a TITAN X. See Table 2 for operator configurations. We also include a weight pre-transformed Winograd [25] for 3x3 conv2d (TVM PT).




Figure 16: ARM A53 end-to-end evaluation of TVM and TFLite.






Figure 17: Relative speedup of all conv2d operators in ResNet-18 and all depthwise conv2d operators in mobilenet. Tested on ARM A53. See Table 2 for the configurations of these operators.

Evaluation

- 3) Embedded GPU
(ARM Mali-T860MP4 GPU on Firefly-RK3399 board):
 - End-to-end performance: TVM outperforms
- 4) FPGA (Vanilla Deep Learning Accelerator-VDLA):
 - TVM outperforms by 40x speedup.
 - The overall performance wasn't estimated since some parts were executed using CPU (gray parts).



Conclusion

- TVM proposed an end-to-end compilation stack for DL across diverse HW backends.
- Graph-level optimization + Operator-level optimization (Auto TVM)
- TVM is evaluated to outperform existing benchmarks on diverse backend platforms.